

SEATauBench: Adapting Tool-Agent-User Evaluation Into Low-Resource Southeast Asian Languages

My Chiffon Nguyen¹, Aulia Adila¹, Saksorn Ruangtanusak^{1,2}, Kittiphat Leesombatwathana^{1,3},
Vissuta Gunawan Lim¹, Patomporn Payoungkhamdee^{1,4}, Samuel Cahyawijaya^{1,5}

¹SEACrowd, ²SCB DataX, ³Chulalongkorn University, ⁴VISTEC, ⁵Cohere
{chiffonng136, auliaadila036, vglim3653}@gmail.com
saksorn.ruangtanusak@data-x.ai, 6534404823@student.chula.ac.th
patomporn.p_s21@vistec.ac.th, samuelcahyawijaya@cohere.com

Abstract

While AI development and evaluation for Southeast Asia (SEA) has grown rapidly, agent capabilities in regional languages are still poorly understood despite its importance to sovereign AI. To fill this gap, we introduce SEATauBench¹, the first agent-focused evaluation framework for SEA sovereign AI. SEATauBench adapts τ^2 -Bench to five languages—Mandarin, Vietnamese, Thai, Indonesian, and Filipino—and evaluates agents across progressively localized settings that vary the language of user-agent interaction, tool specifications, and task domains. Across three recent models, we find that English agent capabilities transfer reasonably well when only the conversation language changes, but quality and robustness degrade sharply as more task contexts are localized, with the largest losses in full domain adaptation. We also the limits of English-only agent assessment for measuring agent capabilities in SEA languages. More broadly, SEATauBench provides a diagnostic benchmark and reusable adaptation pipeline for building reliable multilingual agents for linguistically diverse regions. Data and code can be accessed at github.com/SEACrowd/SEATauBench.

1 Introduction

Sovereign artificial intelligence (AI) has become critical for nations seeking to maintain autonomy in their digital futures (Chae et al., 2025), including Southeast Asia. As articulated in Mushkani et al. (2025); Barasa et al. (2026), sovereign AI encompasses not only technological self-reliance but also cultural and linguistic relevance, a vital dimension for the more than 700 million people whose linguistic diversity is poorly represented by English-centric development and evaluation (Bhandari and Modi, 2026; Putra, 2024).

¹SEATauBench is pronounced "si-tau-bench", similar to the Filipino word for string beans, "sitaw".




Figure 1: SEATauBench exposes a critical English-SEA agentic capability gap in existing proprietary and open-source LLMs across progressively localized evaluation scenarios. This evidence exposes the unreliability of existing English-centric benchmark to reflect the actual capabilities of LLMs for sovereign AI adoption.

SEA-specific language evaluation has expanded rapidly, with SEA-Exam and SEA-Bench (Liu et al., 2025; Zhang et al., 2024; Nguyen et al., 2024), SEA-VL (Cahyawijaya et al., 2025a, 2026), SEACrowd (Lovenia et al., 2024), SEA-HELM (Susanto et al., 2025), NusaCrowd (Cahyawijaya et al., 2023a), NusaWrites (Cahyawijaya et al., 2023b), and NusaX (Winata et al., 2023) establishing valuable foundations for measuring regional language understanding, cultural knowledge, reasoning, safety, and multimodal capabilities. Nevertheless, these benchmarks primarily

assess static model behaviors, with limited work evaluating whether agents can complete multi-turn, tool-mediated tasks in regional languages, the capabilities required when more and more AI systems operate in real-world deployments like service, commerce and travel (Budzianowski et al., 2018; Eric et al., 2020; Zang et al., 2020; Han et al., 2021; Ye et al., 2022).

To address this gap, we introduce SEATauBench, the first agent-focused evaluation framework for SEA. SEATauBench adapts τ^2 -Bench (Yao et al., 2024; Barres et al., 2025; Shi et al., 2026; Ray et al., 2026) into five target languages L2 (Mandarin Chinese, Vietnamese, Thai, Indonesian, Filipino) and evaluates agents across three progressively localized settings: (1) **L2 Interaction**, isolating linguistic capability in user-agent conversation; (2) **L2 Tool**, testing the ability to use tools with non-English tool specifications; and (3) **L2 Domain**, evaluating performance when all task contexts are in L2 (Section 3.3). To translate different interfaces an AI agent interacts with, without breaking the execution of τ^2 -Bench, we develop a structured, non-breaking translation pipeline (Section 3.2).

Across three recent agent models, we find that English agent capabilities transfer reasonably well when agents only need to respond in target languages, but quality and robustness degrade sharply once tools, policies, and task contexts are progressively provided in SEA languages (Section 5.1). These results expose a gap between the growth of SEA evaluation resources and the readiness of current agents for sovereign AI deployment, establishing SEATauBench as a diagnostic benchmark for building reliable multilingual agents for the region.

2 Related Works

2.1 Agent Evaluation

Task-oriented dialogue benchmarks such as MultiWOZ (Budzianowski et al., 2018; Eric et al., 2020; Zang et al., 2020; Han et al., 2021; Ye et al., 2022) and MASSIVE (FitzGerald et al., 2023) established evaluation for goal-directed dialogue and multilingual intent-slot understanding. Recent tool-use benchmarks, including ToolEval (Qin et al., 2024) and BFCL (Patil et al., 2025), extend this direction to function calling and API use. However, deployed service agents must further sustain multi-turn interactions, follow domain policies, and update external states. τ^2 -Bench (Yao et al., 2024; Barres et al., 2025; Shi et al., 2026; Ray et al., 2026)

address this through simulated user-agent-tool environments, scored by task completion metrics pass^1 and pass^3 . Related benchmarks further assess realistic agent behavior in business and professional settings (Huang et al., 2025; Drouin et al., 2024; Boisvert et al., 2024; Xu et al., 2026; Patwardhan et al., 2026).

Despite these advances, existing agent benchmarks remain largely centered on high-resource languages, limiting their applicability to sovereign AI settings. SEATauBench addresses this gap by extending τ^2 -Bench to SEA languages through localized conversational and tool-use scenarios. To our knowledge, it is the first multilingual benchmark to preserve the task-oriented evaluation framework of τ^2 -Bench, contributing towards reliable assessment of agent capabilities in multilingual contexts.

2.2 Multilingual Evaluation

Multilingual evaluation has mainly targeted understanding, reasoning, translation, and multimodal comprehension through benchmarks such as mMMLU (Hendrycks et al., 2021; OpenAI, 2024), GlobalMMLU (Singh et al., 2025), GlotEval (Luo et al., 2025), CVQA (Romero et al., 2024), and AyaVisionBench (CohereLabs, 2025). For South-east Asian languages, resources such as NusaX (Winata et al., 2023), NusaWrites (Cahyawijaya et al., 2023b), NusaDialogue (Purwarianti et al., 2025), IndoTOD (Kautsar et al., 2023), SEACrowd (Lovenia et al., 2024), StingrayBench (Cahyawijaya et al., 2025b), SEAExam/SEABench (Liu et al., 2025), SEA-HELM (Susanto et al., 2025), SEA-VL (Cahyawijaya et al., 2025a), and SEA-VQA (Urailertprasert et al., 2024) expand evaluation to local linguistic, cultural, reasoning, safety, and multimodal contexts. However, these benchmarks largely focus on static evaluation rather than interactive tool use or multi-turn task completion.

The closest work is MASSIVE-Agents (Kulkarni et al., 2025), which evaluates function calling across 52 languages using BFCL, but remains limited to function selection and argument prediction. SEATauBench goes further by evaluating full user-agent-tool interaction, requiring agents to interact over multiple turns, follow instructions, use tools, and complete realistic tasks in SEA languages. It therefore bridges multilingual function calling and interactive agent evaluation, addressing a gap left by existing SEA multilingual and agent evaluation.




Figure 2: Overview of the automated translation pipeline for generating multilingual τ^2 -Bench artifacts (Section 3.2). We provide more details about the pipeline and the resulting translated artifacts in Appendix A.

3 Adapting τ^2 -Bench to SEATauBench

3.1 Background

We construct SEATauBench by extending the English tool-agent-user benchmark τ^2 -Bench (Barres et al., 2025) to multilingual settings. This adaptation must account for two coupled agent-facing surfaces: **interaction content** (task definitions, domain policies and workflows, and structured databases) and **executable interfaces** (tool schemas and, in telecom, return messages).

At runtime, the agent must read policies, reason about tasks, call tools, inspect outputs, and talk to a simulated user in the target language. The benchmark therefore needs a localized interface whose visible text is translated, while the execution layer still preserves canonical English values. Our pipeline aligns these representations to preserve task semantics and metric comparability.

Section 3.2 describes how we translate and construct L2 artifacts, and Section 3.3 uses those artifacts to define controlled evaluation scenarios with increasing levels of L2 adaptation.

3.2 SEATauBench L2 Adaptation Pipeline

Figure 2 summarizes our two-phase adaptation pipeline: **offline translation**, which materializes language-specific assets, and **runtime localization**, which patches the environment observed by the

agent. The full details are provided in Appendix A.

We adapt three domains from τ^2 -Bench (retail, airline, and telecom) into five L2: Vietnamese (vi), Indonesian (id), Thai (th), Filipino (tl), and Mandarin Chinese (zh).²

Offline Translation. We translate static domain assets to construct L2 artifacts. It extracts natural-language spans from tasks, policies, databases, tool docstrings, schemas, and tool-return templates, while masking executable tokens such as IDs, status values, tool names, and structural markers. It first translates schema literals to establish a glossary, preventing the same executable value from being rendered inconsistently. The pipeline then writes outputs with format-specific writers and records per-language manifests with model metadata and source-file SHA-256 fingerprints. Appendix A.5 reports the resulting artifact statistics.

Runtime Localization. The runtime localization phase handles dynamic content exposed during inference. First, it localizes the tool schema shown to the agent: descriptions, enum choices, and examples are rendered in L2 so the scenario tests target-language tool use rather than English schema reading, while the underlying implementation remains unchanged. Second, it preserves executability by

²We do not include other SEA languages such as Malay, Lao, or Cambodian due to limited annotator capacity.

Scenario	L2 Tools	L2 Convo	L2 DB & Policy
(S1) English Only	✗	✗	✗
(S2) L2 Interaction	✗	✓	✗
(S3) L2 Tool	✓	✗	✗
(S4) L2 Domain	✓	✓	✓

Table 1: SEATauBench evaluation scenarios across English and five SEA languages. S2 and S3 isolate L2 dialogue and L2 tool docstring respectively, while S4 L2 domain uses all translated contexts.

normalizing localized arguments back to canonical English values before tool calls; otherwise, correct L2 arguments could fail only because the original tools expect English literals. After execution, tool responses are localized back into L2, and final payloads are canonicalized again before scoring, keeping the interaction monolingual for the agent while making metrics comparable across languages and with the English benchmark. Together, offline translation and runtime localization let us vary which benchmark surfaces are exposed in L2, forming the scenario design.

Human Manual Review. For each target language, a native speaker—either an author or a recruited reviewer—reviews the machine-translated artifacts. They read every translated prose document in full (domain policy and agent/user instructions) and a sample of 100 segments each from the database and the task contexts. Reviews are carried out in a per-language Excel workbook with one sheet per artifact (see details in Appendix ??), where annotators add a corrected translation and notes (both optional). If human corrected values are present, we use them in experiments; otherwise, we consider machine translated versions usable. We find that our pipeline ingests 91% of machine-translated artifacts, showing that they can be trusted even if used by themselves. Since we do not employ more than one reviewer per language, we do not calculate inter-annotator agreement scores.

3.3 Evaluation Scenarios

SEATauBench has four scenarios with increasing level of L2 adaptation (summary in Table 1).

(S1) English Only. The original English benchmark, to establish the baseline performance for tested agents.

(S2) L2 Interaction. The simulated user and agent converse in L2, while tools, policies, databases,

and task contexts remain in English.

(S3) L2 Tool. Tool schemas are rewritten in L2 while conversation and domain context remain in English. We present translated tool schemas in two formats: single-L2, where each schema uses one target language, and mixed-L2, where schemas combine an increasing number of target languages. **(S4) L2 Domain.** Agents operate in a complete L2 setting where dialogue, tool schemas, policies, task descriptions, and agent-visible database are translated while task semantics remain fixed. This scenario uses all translated artifacts and runtime localization described in Section 3.2.

4 Experimental Setup

Metrics. Following τ^2 -Bench (Barres et al., 2025), we report two task metrics. First, $\text{pass}@1$ measures quality as the mean success rate of independent trials reaching the expected final state, i.e., $\text{pass}@1 = \text{pass}^1 = \mathbb{E}[r]$. Building on existing metrics, we also derive a robustness measure, $\rho \in [0, 1]$, defined as

$$\rho^q = \frac{\text{pass}^q}{\text{pass}@1} \quad (1)$$

where ρ^q denotes the probability that all q independent trials succeed, averaged across tasks. $\rho^q = 1$ indicates that the agent performs consistently across multiple runs, while lower ρ^q indicates that the agent may solve a task once but fail to do so reliably across repeated trials.

Task and Domain. For high-quality evaluation, we use latest versions of retail and airline domains (Cuadron et al., 2026), which incorporate several corrections and refinements over prior versions (Yao et al., 2024; Barres et al., 2025). For the telecom domain, we follow the original τ^2 -Bench, as its tasks were not updated in the most recent release. We leave exploration of the banking domain (Shi et al., 2026) and the voice modality (Ray et al., 2026) introduced in τ^3 -bench to future work.

Hyperparameters and Models. For all evaluations, we use $q = 3$, resulting in the robustness metric ρ^3 , and Qwen3-235B-A22B-Inst³ (Qwen-Team, 2025) as the user LLM. For the natural-language assertion judge, we use GPT-4.1, following the implementations provided in τ^2 -Bench. We evaluate three recent LLM agents spanning three model families: GPT-5-Mini⁴ (Singh et al., 2026) as a proprietary model, Qwen3-235B-A22B-Instruct-2507

³<https://huggingface.co/Qwen/Qwen3-235B-A22B-Inst-A22B-Instruct-2507>

⁴<https://developers.openai.com/api/docs/models/gpt-5-mini>




Figure 3: Evaluation results on SEATauBench. SEATauBench reveals consistent degradation as the setting becomes less English-centric on all languages in both pass@1 and ρ^3 metrics, with most severe drops in low-resource languages with non-Latin scripts. All scenario results are reported with three models, except for L2 Tool, where only GPT-5-mini and Qwen3-235B-A22B-Inst are reported.

as a representative open-source model, and Kimi-K2.5 (Kimi Team et al., 2026) as a larger-scale open-source model.

For each model, we use the default hyperparameters defined in each provider (See Appendix C). We avoid any model-specific hyperparameter tuning to better reflect the out-of-the-box tool-agent-user interaction capability when adopted to the specified scenario and L2.

5 Results

We present two main results. First, we examine how quality and robustness change as more components of the English benchmark are converted to L2 (Section Section 5.1). Second, we test whether these trends are consistent across agent models and task domains (Section Section 5.2). Together, the results show that English-centric evaluations overestimate multilingual agent reliability, and that the size of this gap depends on both the target language and the structure of the task environment.

5.1 Performance Degrades in L2 Languages

Figure 3 demonstrates a consistent decline in model performance as the evaluation setting becomes progressively less English-centric. Both pass@1 and ρ^3 are strongest in English settings, but perfor-

mance drops when the user-agent conversation is conducted in an L2, drops further when tool interfaces are localized, and typically reaches its lowest levels when the domain context is also translated. This progression indicates that multilingual agentic competence is not only a matter of understanding non-English user utterances. Agents must also use localized tool schemas, entities, and ground decisions in domain policies expressed in L2.

Although all L2 settings degrade relative to English, the severity and form of degradation differ by language. In L2 Domain, Thai and Filipino show the largest pass@1 declines, falling to roughly the low-0.4 range. Vietnamese and Chinese remain closer to the mid-0.5 range, while Indonesian shows a partial robustness recovery despite lower pass@1. The sharper drop in Thai may reflect challenges associated with a distinct writing system and lower-resource coverage. However, the comparable drop in Filipino shows that localization failures are not limited to non-Latin scripts: even Latin-script languages can expose weaknesses when agents must reason over fully localized dialogue, tools, schemas, and domain context.

Quality and robustness diverge. Figure 3 also shows that pass@1 and ρ^3 do not always move together. In Indonesian, Filipino, Chinese, and to a




Figure 4: Quality and robustness of varying agent models across languages and domains. **(left)** All agents yield consistently higher $\text{pass}@1$ in English and Chinese indicating a quality bias toward high-resource languages. **(center)** A similar trend is observed on ρ^3 , with negative correlation of ρ^3 between English and L2 languages especially Thai. **(right)** The trend of $\text{pass}@1$ holds across most domains and models with some exceptions as described in §5.2.

lesser extent Vietnamese, $\text{pass}@1$ drops in L2 Domain while ρ^3 remains comparatively stable or rebounds relative to L2 Tool. This means that agents may solve fewer tasks overall, but when they do solve a task, their successful behavior is relatively reproducible in some L2 settings. Thai is the clearest exception: both $\text{pass}@1$ and ρ^3 decline in L2 Domain. This suggests that some L2 conditions weaken not only average task success but also consistency across repeated trials.

5.2 Does the Trend Hold across Agents?

Overall, model performance differs across metrics, as shown in Fig. 4. For $\text{pass}@1$, Kimi K2.5 performs best in most languages, including English, Vietnamese, Filipino, and Chinese, while GPT-5 Mini leads in Indonesian and Thai, indicating that model advantages do not transfer uniformly across SEA languages. Qwen3-235B-A22B-Inst consistently has the lowest $\text{pass}@1$, suggesting weaker single-run task completion. However, ρ^3 reveals a different pattern: despite its lower $\text{pass}@1$, Qwen3-235B-A22B-Inst often shows stronger robustness in lower-resource L2 settings. This robustness advantage weakens in higher-resource languages such as English and Chinese, where Kimi K2.5 and GPT-5 Mini remain competitive or stronger.

The domain-level results show that the English–non-English gap varies across domains and models. For GPT-5 Mini and Kimi K2.5, $\text{pass}@1$ generally follows Telecom > Airline > Retail in both English and non-English settings, suggesting that Retail is the most challenging domain due to its policies, item-level constraints, and state updates. Qwen3-235B-A22B-Inst Instruct shows a different pattern: Retail performs best in English but

becomes weakest in non-English settings, indicating that domain-specific strengths in English may not transfer under localization. The ρ^3 results further show that localization reduces robustness, with English settings generally achieving higher consistency than non-English settings across domains. However, this degradation is not uniform. For Qwen3-235B-A22B-Inst, agent performance follows the expected non-English decline in Telecom and Retail domains, while shows a relative robustness spike for Airline. This suggests that robustness under localization depends on both task domain and the agent model behavior.

6 Analysis and Discussion

In this section, we examine error patterns arising from conversational user–agent interactions in L2 Interaction and L2 Domain scenarios (Section 6.1) and the effects of having tools in multiple languages (Section 6.2). Furthermore, we show that agent performance has little to no association with language correctness (Section 6.3) and that English performance is an unreliable proxy for L2 performance with SEA languages (Section 6.4).

6.1 Error Analysis

To further examine the downstream performance metrics reported in the previous section, we conduct a qualitative error analysis based on the LLM-as-a-Judge framework proposed in Shi et al. (2026). Following Barres et al. (2025), we categorize task simulation errors into two types: *Critical* error, in which the agent deviates from the user’s intent or commits irrecoverable mistakes, and *Benign* error, which do not prevent successful task completion. For early results, we analyze only scenarios that af-



Figure 5: Error categorization for simulated user and agent, across SEA language in the L2 Interaction and L2 Domain scenarios. Both agents and users tend to have higher critical errors in L2 than English, while having a similar benign error in comparison to English. The errors are reported by judge DeepSeek-V4-Flash, for simulated user Qwen3-235B-A22B-Inst and agent GPT-5-mini.

fect both behavioral changes on both simulated user and agent, i.e., L2 Interaction and L2 Domain scenarios, since L2 Tool only translates tool specifications while keeping all tool-agent-user interactions in English. Specifically, we employ DeepSeek-V4-Flash as the oracle LLM on simulations done by GPT-5-mini as the agent and Qwen3-235B-A22B-Inst as the simulated user.

L2 Interaction. Agent performance degrades when interacting with users in L2. As shown in Figure 5, critical errors appear in around 40% of total simulations across SEA languages, with Thai being the most severe (nearly 50%). Moreover, when the user model must communicate only in SEA languages, its behavior becomes less stable and less consistent with the intended simulation setup, with the most severe degradation in Filipino.

L2 Domain. Agent performance degradation is more pronounced in the L2 domain adaptation setting. There are slightly more critical errors for agents here than L2 interaction, with clearest trends in Thai, Filipino, and Indonesian (Figure 5). This finding suggests that operating in complete L2 environments further impairs the agent’s ability to complete tasks effectively.

Simulated User. A similar trend is observed on the simulated user, where critical errors consistently account for around 20% of all errors in




Figure 6: Pass@1 and ρ^3 as the number of languages used for translated tools increase from 1 (English) to 5. Performance for two agents (GPT-5-Mini & Qwen3-235B-A22B-Inst) slightly drops when a second language is added, but it plateaus for more languages.

both analyzed scenarios. Our main finding that agent performance drops as L2 adaptation increases across four scenarios (Figure 3) might be compromised by the low reliability of the simulated user. This account previously appeared in Barres et al. (2025) for English and Seshadri et al. (2026) for dialectical English.

6.2 Mixed language tool causes modest drop in task performance

Beyond adaptation using a single L2, we also investigate how agents handle tools in multiple languages, which is increasingly relevant as tools for




Figure 7: Only 1.4% of the variation in task performance metric pass^3 is explained by language use capability of tested agents, suggesting a very weak linear relationship. Each dot represents a run for one scenario $\times$ domain $\times$ language setting.

agents are developed in different world regions. We do so by extending L2 Tool experiment and explore five different mixed language settings by gradually adding more languages to the mix for tool specification: Mix-2 uses English and Thai, Mix-3 adds Vietnamese, Mix-4 adds Indonesian, and Mix-5 adds Chinese (for more details, see Appendix B).

Figure 6 shows that mixed language tool causes only a modest drop. $\text{Pass}@1$ drops from 0.68 for English-only average to 0.55 across all mixed-language settings (0.13 decrease). ρ^3 shows a similar trend of degradation (0.11 decrease). These results indicate that agents with strong tool-using performance in English might use tools slightly less effectively when tools are multilingual; however, increasing the language diversity in tool specification has little to no impact to the overall agentic tool-using capability.

6.3 Language use is uncorrelated with performance

We also measure whether the simulated user and agent operated across tasks in the correct L2, or *language correctness*. It is calculated as the fraction of eligible turns whose detected language (by fast-Text) matches the expected L2 for a single run. We exclude trajectories with system errors and other errors unrelated to task execution.

Figure 7 shows that run-level language correctness for agent has little association with task performance: a linear fit explains only $R^2 = 0.014$ of the variance in pass^3 . Several runs achieve near-

perfect language correctness while spanning a wide range of pass^3 scores, especially in L2 Tool and L2 Domain. Conversely, L2 Interaction contains runs with visibly lower language correctness but still moderate or high task success. This suggests that, at the aggregate run level, task failures are not primarily explained by whether the agent remains in the target language. For language drift, we also measure how often the agent produces text in the expected L2 (Figure 9), where in a conversation off-target text appears (Figure 10), and how much of the off-target text is specifically English (Figure 11).

6.4 Is English a reliable proxy for L2 in SEA?

The τ^2 -Bench evaluation framework, while effective for measuring agent capabilities, is resource-intensive and computationally expensive. Extending to SEATauBench with an additional language dimension adds to these costs. To address this challenge, we examine correlation between quality and robustness metrics across SEATauBench languages to determine whether performance in one language could serve as a proxy for other languages.

As illustrated in Figure 8, agent performance in English has high predictive power of L2 performance for $\text{pass}@1$, with correlation above 0.9. However, there is more variability in ρ^3 , where the highest correlation is in Chinese (0.88) and the lowest in Thai (0.49). Therefore, English performance is not a reliable representation for SEA language performance when we factor in consistency. Filipino (TL), however, shows the strongest correlation with other SEA languages, with the lowest score being 0.91 for $\text{pass}@1$ and 0.85 for ρ^3 . Based on these findings, we recommend using the Filipino subset of SEATauBench as a proxy to estimate performance across the remaining four languages, for a more efficient yet representative evaluations of AI agents in the SEA region.

7 Conclusion

SEATauBench establishes the first agent-focused framework for SEA linguistic diversity, revealing that English agentic capabilities transfer effectively to a superficial L2 interaction but degrade significantly with the inclusion of non-English contexts, due to both agent limitations and reduced user capacity. By exposing fundamental gaps in English-only evaluation methodologies as a proxy of agents capability in SEA languages, SEATauBenchpro-




Figure 8: Correlation of pass@1 (upper triangle) and ρ^3 (lower triangle) across English and SEATauBench languages. Correlation scores are aggregated by four scenarios, three agent models, and three task domains.

vides essential diagnostic tools for sovereign AI development, establishing a rigorous foundation for future research in sovereign agentic AI solutions, ultimately supporting autonomous AI across SEA diverse linguistic communities.

8 Limitation

This work presents a systematic framework for extending agentic benchmarks to evaluate multilingual conversational capabilities with tool access in Southeast Asian (SEA) languages, specifically Vietnamese, Indonesian, Thai, Filipino, and Chinese, across the retail, airline, and telecom domains. Although the proposed framework can, in principle, be extended to additional languages, the empirical findings and their implications are confined to the SEA language setting explored in this study.

Another limitation is the benchmark adaptation methodology, which is specifically designed around the τ -Bench evaluation paradigm (Yao et al., 2024; Barres et al., 2025). While the proposed evaluation scenarios may be applicable more broadly, the underlying design constraints may limit the generalizability of the methodology beyond τ -Bench frameworks. In addition, the current analysis is limited to English and five Southeast Asian languages, consequently, the observed findings may not generalize to other SEA languages that are not covered in this work.

9 Ethical Statement

We conducted this research with careful attention to ethical considerations throughout the development process. In developing of the translation pipeline and the evaluation framework, we ensure that there is no external sensitive information is exposed throughout the process and the resulting artifacts. We recognize our responsibility to ensure these systems are fair, respectful of different cultures, and beneficial to all users. Our approach focuses on being transparent about what the agents can and cannot do, actively working to identify and address potential biases, and taking steps to prevent misuse. We handle all research data responsibly, protecting privacy while respecting linguistic diversity. This work aims to contribute to AI that serves people equitably across languages and cultures, and we are committed to being accountable for how these technologies impact society.

Acknowledgment

This work was supported by the ThaiLLM collaboration, funded by the Digital Economy and Society (DE) Development Fund of the Ministry of Digital Economy and Society, Thailand. We thank Kian Kyars for providing OpenAI API access for our experiments, and SCBX R&D for providing the necessary Google Cloud Vertex AI resources.

References

- Hilda Barasa, PeiChin Tay, Keegan McBride, Alexander Iosad, and Jakob Mökander. 2026. [Sovereignty in the Age of AI: Strategic Choices, Structural Dependencies](#). Accessed: 22 May 2026.
- Victor Barres, Honghua Dong, Soham Ray, Xujie Si, and Karthik Narasimhan. 2025. [\$\tau^2\$ -bench: Evaluating conversational agents in a dual-control environment](#). Preprint, arXiv:2506.07982.
- Manik Bhandari and Gaurav Modi. 2026. [Why sovereign artificial intelligence is imperative in south-east asia](#). Accessed: 22 May 2026.
- Léo Boisvert, Megh Thakkar, Maxime Gasse, Massimo Caccia, Thibault Le Sellier De Chezelles, Quentin Cappart, Nicolas Chapados, Alexandre Lacoste, and Alexandre Drouin. 2024. [Workarena++: Towards compositional planning and reasoning-based common knowledge work tasks](#). In *Advances in Neural Information Processing Systems*, volume 37, pages 5996–6051. Curran Associates, Inc.
- Paweł Budzianowski, Tsung-Hsien Wen, Bo-Hsiang Tseng, Iñigo Casanueva, Stefan Ultes, Osman Ra-

- madan, and Milica Gašić. 2018. [MultiWOZ - a large-scale multi-domain Wizard-of-Oz dataset for task-oriented dialogue modelling](#). In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 5016–5026, Brussels, Belgium. Association for Computational Linguistics.
- Samuel Cahyawijaya, Peerat Limkonchotiawat, Tack Hwa Wong, Hitesh Laxmichand Patel, Amit Agarwal, Manuel Antonio Rufino, Carlos Rafael Catalan, Muhammad Reza Qorib, Vicky Feliren, Holy Lovenia, Aye Hninn Khine, Frederikus Hudi, David Anugraha, Alham Fikri Aji, Romrawin Chumpu, Viet-Thanh Pham, Minghan Wang, Mohamed Fazli Imam, Ruochen Zhang, and 29 others. 2026. [Anthropogenic regional adaptation in multimodal vision-language model](#). *Preprint*, arXiv:2604.11490.
- Samuel Cahyawijaya, Holy Lovenia, Alham Fikri Aji, Genta Winata, Bryan Wilie, Fajri Koto, Rahmad Mahendra, Christian Wibisono, Ade Romadhony, Karissa Vincentio, Jennifer Santoso, David Moeljadi, Cahya Wirawan, Frederikus Hudi, Muhammad Satrio Wicaksono, Ivan Parmonangan, Ika Alfina, Ilham Firdausi Putra, Samsul Rahmadani, and 29 others. 2023a. [NusaCrowd: Open source initiative for Indonesian NLP resources](#). In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 13745–13818, Toronto, Canada. Association for Computational Linguistics.
- Samuel Cahyawijaya, Holy Lovenia, Fajri Koto, Dea Adhista, Emmanuel Dave, Sarah Oktavianti, Salsabil Akbar, Jhonson Lee, Nur Shadieq, Tjeng Wawan Cenggoro, Hanung Linuwih, Bryan Wilie, Galih Muridan, Genta Winata, David Moeljadi, Alham Fikri Aji, Ayu Purwarianti, and Pascale Fung. 2023b. [NusaWrites: Constructing high-quality corpora for underrepresented and extremely low-resource languages](#). In *Proceedings of the 13th International Joint Conference on Natural Language Processing and the 3rd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 921–945, Nusa Dua, Bali. Association for Computational Linguistics.
- Samuel Cahyawijaya, Holy Lovenia, Joel Ruben Antony Moniz, Tack Hwa Wong, Mohammad Rifqi Farhan-syah, Thant Thiri Maung, Frederikus Hudi, David Anugraha, Muhammad Ravi Shulthan Habibi, Muhammad Reza Qorib, Amit Agarwal, Joseph Marvin Imperial, Hitesh Laxmichand Patel, Vicky Feliren, Bahrul Ilmi Nasution, Manuel Antonio Rufino, Genta Indra Winata, Rian Adam Rajagede, Carlos Rafael Catalan, and 73 others. 2025a. [Crowd-source, crawl, or generate? creating SEA-VL, a multicultural vision-language dataset for Southeast Asia](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 18685–18717, Vienna, Austria. Association for Computational Linguistics.
- Samuel Cahyawijaya, Ruochen Zhang, Jan Christian Blaise Cruz, Holy Lovenia, Elisa Gilbert, Hiroki Nomoto, and Alham Fikri Aji. 2025b. [Thank you, stingray: Multilingual large language models can not \(yet\) disambiguate cross-lingual word senses](#). In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 3228–3250, Albuquerque, New Mexico. Association for Computational Linguistics.
- Kyubyung Chae, Gihoon Kim, Gyuseong Lee, Taesup Kim, Jaemin Lee, and Heejin Kim. 2025. [Assessing socio-cultural alignment and technical safety of sovereign LLMs](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 10579–10600, Suzhou, China. Association for Computational Linguistics.
- CohereLabs. 2025. [AyaVisionBench](#). <https://huggingface.co/datasets/CohereLabs/AyaVisionBench>. Hugging Face dataset.
- Alejandro Cuadron, Pengfei Yu, Yang Liu, and Arpit Gupta. 2026. [SABER: Small actions, big errors — safeguarding mutating steps in LLM agents](#). In *ICLR 2026 Workshop on Memory for LLM-Based Agentic Systems*.
- Alexandre Drouin, Maxime Gasse, Massimo Caccia, Issam H. Laradji, Manuel Del Verme, Tom Marty, Léo Boisvert, Megh Thakkar, Quentin Cappart, David Vazquez, Nicolas Chapados, and Alexandre Lacoste. 2024. [Workarena: How capable are web agents at solving common knowledge work tasks?](#) *Preprint*, arXiv:2403.07718.
- Mihail Eric, Rahul Goel, Shachi Paul, Abhishek Sethi, Sanchit Agarwal, Shuyang Gao, Adarsh Kumar, Anuj Goyal, Peter Ku, and Dilek Hakkani-Tur. 2020. [MultiWOZ 2.1: A consolidated multi-domain dialogue dataset with state corrections and state tracking base-lines](#). In *Proceedings of the Twelfth Language Resources and Evaluation Conference*, pages 422–428, Marseille, France. European Language Resources Association.
- Jack FitzGerald, Christopher Hench, Charith Peris, Scott Mackie, Kay Rottmann, Ana Sanchez, Aaron Nash, Liam Urbach, Vishesh Kakarala, Richa Singh, Swetha Ranganath, Laurie Crist, Misha Britan, Wouter Leeuwis, Gokhan Tur, and Prem Natara-jan. 2023. [MASSIVE: A 1M-example multilingual natural language understanding dataset with 51 typologically-diverse languages](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4277–4302, Toronto, Canada. Association for Computational Linguistics.
- Ting Han, Ximing Liu, Ryuichi Takanabu, Yixin Lian, Chongxuan Huang, Dazhen Wan, Wei Peng, and Minlie Huang. 2021. [Multiwoz 2.3: A multi-domain task-oriented dialogue dataset enhanced with annotation corrections and co-reference annotation](#). In *Natural Language Processing and Chinese Computing*, pages 206–218, Cham. Springer International Publishing.

- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021. [Measuring massive multitask language understanding](#). In *International Conference on Learning Representations*.
- Kung-Hsiang Huang, Akshara Prabhakar, Sidharth Dhawan, Yixin Mao, Huan Wang, Silvio Savarese, Caiming Xiong, Philippe Laban, and Chien-Sheng Wu. 2025. [CRMarena: Understanding the capacity of LLM agents to perform professional CRM tasks in realistic environments](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3830–3850, Albuquerque, New Mexico. Association for Computational Linguistics.
- Muhammad Kautsar, Rahmah Nurdini, Samuel Cahyawijaya, Genta Winata, and Ayu Purwarianti. 2023. [InDoToD: A multi-domain Indonesian benchmark for end-to-end task-oriented dialogue systems](#). In *Proceedings of the First Workshop in South East Asian Language Processing*, pages 85–99, Nusa Dua, Bali, Indonesia. Association for Computational Linguistics.
- Kimi Team, Tongtong Bai, Yifan Bai, Yiping Bao, S. H. Cai, Yuan Cao, Y. Charles, H. S. Che, Cheng Chen, Guanduo Chen, Huarong Chen, Jia Chen, Jiahao Chen, Jianlong Chen, Jun Chen, Kefan Chen, Liang Chen, Ruijue Chen, Xinhao Chen, and 307 others. 2026. [Kimi k2.5: Visual agentic intelligence](#). *Preprint*, arXiv:2602.02276.
- Mayank Kulkarni, Vittorio Mazzia, Judith Gaspers, Chris Hench, and Jack FitzGerald. 2025. [MASSIVE-agents: A benchmark for multilingual function-calling in 52 languages](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 20193–20215.
- Chaoqun Liu, Wenxuan Zhang, Jiahao Ying, Mahani Aljunied, Anh Tuan Luu, and Lidong Bing. 2025. [SeaExam and SeaBench: Benchmarking LLMs with local multilingual questions in Southeast Asia](#). In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 6134–6151, Albuquerque, New Mexico. Association for Computational Linguistics.
- Holy Lovenia, Rahmad Mahendra, Salsabil Maulana Akbar, Lester James V. Miranda, Jennifer Santoso, Elyanah Aco, Akhdan Fadhillah, Jonibek Mansurov, Joseph Marvin Imperial, Onno P. Kampman, Joel Ruben Antony Moniz, Muhammad Ravi Shulthan Habibi, Frederikus Hudi, Railey Montalan, Ryan Ignatius, Joanito Agili Lopo, William Nixon, Börje F. Karlsson, James Jaya, and 42 others. 2024. [SEACrowd: A multilingual multimodal data hub and benchmark suite for Southeast Asian languages](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 5155–5203, Miami, Florida, USA. Association for Computational Linguistics.
- Hengyu Luo, Zihao Li, Joseph Attieh, Sawal Devkota, Ona de Gibert, Xu Huang, Shaoxiong Ji, Peiqin Lin, Bhavani Sai Praneeth Varma Mantina, Ananda Sreenidhi, Raúl Vázquez, Mengjie Wang, Samea Yusufi, Fei Yuan, and Jörg Tiedemann. 2025. [GlotEval: A test suite for massively multilingual evaluation of large language models](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 602–614, Suzhou, China. Association for Computational Linguistics.
- Rashid Mushkani, Hugo Berard, Allison Cohen, and Shin Koseki. 2025. [Position: The right to AI](#). In *Forty-second International Conference on Machine Learning Position Paper Track*.
- Xuan-Phi Nguyen, Wenxuan Zhang, Xin Li, Mahani Aljunied, Zhiqiang Hu, Chenhui Shen, Yew Ken Chia, Xingxuan Li, Qingyu Tan Jianyu Wang, Liying Cheng, Guanzheng Chen, Yue Deng, Sen Yang, Chaoqun Liu, Hang Zhang, and Lidong Bing. 2024. [Seallms - large language models for southeast asia](#).
- OpenAI. 2024. [MMMLU: Multilingual massive multitask language understanding](#). <https://huggingface.co/datasets/openai/MMMLU>. Hugging Face dataset.
- Shishir G. Patil, Huanzhi Mao, Charlie Cheng-Jie Ji, Fanjia Yan, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. 2025. [The berkeley function calling leaderboard \(BFCL\): From tool use to agentic evaluation of large language models](#). In *Proceedings of the 42nd International Conference on Machine Learning*.
- Tejal Patwardhan, Rachel Dias, Elizabeth Proehl, Grace Kim, Michele Wang, Olivia Watkins, Simon Posada Fishman, Marwan Aljubei, Phoebe Thacker, Lorraine Fauconnet, Natalie S. Kim, Samuel Misereindino, Gildas Chabot, David Li, Patrick Chao, Michael Sharman, Alexandra Barr, Amelia Glaese, and Jerry Tworek. 2026. [GDPval: Evaluating AI model performance on real-world economically valuable tasks](#). In *The Fourteenth International Conference on Learning Representations*.
- Ayu Purwarianti, Dea Adhista, Agung Baptiso, Miftahul Mahfuzh, Yusrina Sabila, Aulia Adila, Samuel Cahyawijaya, and Alham Fikri Aji. 2025. [NusaDialogue: Dialogue summarization and generation for underrepresented and extremely low-resource languages](#). In *Proceedings of the Second Workshop in South East Asian Language Processing*, pages 82–100, Online. Association for Computational Linguistics.
- Bama Andika Putra. 2024. [Governing ai in southeast asia: Asean’s way forward](#). *Frontiers in Artificial Intelligence*, Volume 7 - 2024.
- Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, Sihan Zhao, Lauren Hong, Runchu Tian, Ruobing Xie, Jie Zhou, Mark Gerstein, dahai li,

- Zhiyuan Liu, and Maosong Sun. 2024. [Toolllm: Facilitating large language models to master 16000+ real-world apis](#). In *International Conference on Learning Representations*, volume 2024, pages 9695–9717.
- Qwen-Team. 2025. [Qwen3 technical report](#). *Preprint*, arXiv:2505.09388.
- Soham Ray, Keshav Dhandhanian, Victor Barres, and Karthik Narasimhan. 2026. [\$\tau\$ -voice: Benchmarking full-duplex voice agents on real-world domains](#). *Preprint*, arXiv:2603.13686.
- David Romero, Chenyang Lyu, Haryo Akbarianto Wibowo, Teresa Lynn, Injy Hamed, Aditya Nanda Kishore, Aishik Mandal, Alina Dragonetti, Artem Abzaliev, Atnafu Lambebo Tonja, Bontu Fufa Balcha, Chenxi Whitehouse, Christian Salamea, Dan John Velasco, David Ifeoluwa Adelani, David Le Meur, Emilio Villa-Cueva, Fajri Koto, Fauzan Farooqui, and 56 others. 2024. [Cvqa: Culturally-diverse multilingual visual question answering benchmark](#). *Preprint*, arXiv:2406.05967.
- Preethi Seshadri, Samuel Cahyawijaya, Ayomide Odumakin, Sameer Singh, and Seraphina Goldfarb-Tarrant. 2026. [Lost in simulation: Llm-simulated users are unreliable proxies for human users in agentic evaluations](#). *Preprint*, arXiv:2601.17087.
- Quan Shi, Alexandra Zyttek, Pedram Razavi, Karthik Narasimhan, and Victor Barres. 2026. [\$\tau\$ -knowledge: Evaluating conversational agents over unstructured knowledge](#). *Preprint*, arXiv:2603.04370.
- Aaditya Singh, Adam Fry, Adam Perelman, Adam Tart, Adi Ganesh, Ahmed El-Kishky, Aidan McLaughlin, Aiden Low, AJ Ostrow, Akhila Ananthram, Akshay Nathan, Alan Luo, Alec Helyar, Aleksander Madry, Aleksandr Efremov, Aleksandra Spyra, Alex Baker-Whitcomb, Alex Beutel, Alex Karpenko, and 467 others. 2026. [Openai gpt-5 system card](#). *Preprint*, arXiv:2601.03267.
- Shivalika Singh, Angelika Romanou, Cl  mentine Fourrier, David Ifeoluwa Adelani, Jian Gang Ngui, Daniel Vila-Suero, Peerat Limkonchotiwat, Kelly Marchisio, Wei Qi Leong, Yosephine Susanto, Raymond Ng, Shayne Longpre, Sebastian Ruder, Wei-Yin Ko, Antoine Bosselut, Alice Oh, Andre Martins, Leshem Choshen, Daphne Ippolito, and 4 others. 2025. [Global MMLU: Understanding and addressing cultural and linguistic biases in multilingual evaluation](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 18761–18799, Vienna, Austria. Association for Computational Linguistics.
- Yosephine Susanto, Adithya Venkatadri Hulagadri, Jann Railey Montalan, Jian Gang Ngui, Xianbin Yong, Wei Qi Leong, Hamsawardhini Rengarajan, Peerat Limkonchotiwat, Yifan Mai, and William Chandra Tjhi. 2025. [SEA-HELM: Southeast Asian holistic evaluation of language models](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 12308–12336, Vienna, Austria. Association for Computational Linguistics.
- Norawit Uraileertprasert, Peerat Limkonchotiwat, Supasorn Suwajanakorn, and Sarana Nutanong. 2024. [SEA-VQA: Southeast Asian cultural context dataset for visual question answering](#). In *Proceedings of the 3rd Workshop on Advances in Language and Vision Research (ALVR)*, pages 173–185, Bangkok, Thailand. Association for Computational Linguistics.
- Genta Indra Winata, Alham Fikri Aji, Samuel Cahyawijaya, Rahmad Mahendra, Fajri Koto, Ade Romadhony, Kemal Kurniawan, David Moeljadi, Radityo Eko Prasoj, Pascale Fung, Timothy Baldwin, Jey Han Lau, Rico Sennrich, and Sebastian Ruder. 2023. [NusaX: Multilingual parallel sentiment dataset for 10 Indonesian local languages](#). In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pages 815–834, Dubrovnik, Croatia. Association for Computational Linguistics.
- Frank F. Xu, Yufan Song, Boxuan Li, Yuxuan Tang, Krittanjali Jain, Mengxue Bao, Zora Zhiruo Wang, Xuhui Zhou, Zhitong Guo, Murong Cao, Mingyang Yang, Hao Yang Lu, Amaad Martin, Zhe Su, Leander Melroy Maben, Raj Mehta, Wayne Chi, Lawrence Kenho Jang, Yiqing Xie, and 2 others. 2026. [Theagent-company: Benchmarking LLM agents on consequential real world tasks](#). In *The Thirty-ninth Annual Conference on Neural Information Processing Systems Datasets and Benchmarks Track*.
- Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. 2024. [\$\tau\$ -bench: A benchmark for tool-agent-user interaction in real-world domains](#). *Preprint*, arXiv:2406.12045.
- Fanghua Ye, Jarana Manotumruksa, and Emine Yilmaz. 2022. [MultiWOZ 2.4: A multi-domain task-oriented dialogue dataset with essential annotation corrections to improve state tracking evaluation](#). In *Proceedings of the 23rd Annual Meeting of the Special Interest Group on Discourse and Dialogue*, pages 351–360, Edinburgh, UK. Association for Computational Linguistics.
- Xiaoxue Zang, Abhinav Rastogi, Srinivas Sunkara, Raghav Gupta, Jianguo Zhang, and Jindong Chen. 2020. [MultiWOZ 2.2 : A dialogue dataset with additional annotation corrections and state tracking baselines](#). In *Proceedings of the 2nd Workshop on Natural Language Processing for Conversational AI*, pages 109–117, Online. Association for Computational Linguistics.
- Wenxuan Zhang, Hou Pong Chan, Yiran Zhao, Mahani Aljunied, Jianyu Wang, Chaoqun Liu, Yue Deng, Zhiqiang Hu, Weiwen Xu, Yew Ken Chia, Xin Li, and Lidong Bing. 2024. [Seallms 3: Open foundation and chat multilingual large language models for southeast asian languages](#).

A Machine Translation Pipeline Details

A.1 τ^2 -BenchContext

Each τ^2 -Benchdomain (airline, retail, telecom) exposes the agent to two coupled surfaces. First, `data/tau2/domains/{domain}/` contains the benchmark content seen during an interaction: task definitions in `tasks.json`, domain policies and workflow documents in markdown files `*.md`, and structured databases in `db.json`, `db.toml`, or user-side DB files. Second, `src/tau2/domains/{domain}/` contains the executable tool interface: toolkits in `tools.py` and `user_tools.py`, schema definitions in `data_model.py` and `user_data_model.py`, and in telecom, tool return string templates.

At runtime, the agent must read policy, reason about tasks, call tools, inspect tool outputs, and communicate with a simulated user entirely in the target language.

A.2 Translation Principles

Adapting τ^2 -Bench to non-English languages requires translating the **natural-language content** seen by the language model while leaving untouched the **runtime-canonical tokens** consumed by the tool execution layer: enum literals, entity identifiers, status codes, and tool argument names. Conflating these two layers would corrupt evaluation: a translated product name that reaches the execution engine might trigger a lookup failure, while an untranslated persona description breaks the monolingual interaction we intend to test. We therefore build the pipeline around four principles: **Runtime invariance**. Canonical tokens remain English at execution time. Only the LLM-visible surface is localized.

Terminological consistency. A canonical value should surface with the same target-language realization everywhere it appears.

Format fidelity. Translation is applied only to natural-language leaves. Structural keys, numeric fields, booleans, and program identifiers are never modified.

Bidirectional transparency. Localized values can be mapped back to canonical English forms before execution and before metric computation.

The implementation has two phases: an **offline translation phase** that materializes language-specific assets, and a **runtime localization phase** that patches the environment presented to the agent.

A.3 Phase 1. Offline Translation

The offline phase consumes a domain’s static artifacts, translate with **Gemini 3.1 Flash Lite**, and writes translated outputs to `data/tau2/domains/{domain}/{lang_id}/`.

Step 1. Artifact discovery For each selected domain, the pipeline discovers files from `data/tau2/domains/{domain}/` and `src/tau2/domains/{domain}/`, then assigns each file a processing kind (markdown, json, toml, or python). Refer to Table 2 for the list of artifacts and their role in translation.

Step 2. Segment extraction with path-sensitive rules. Each file is reduced to minimal translatable segments together with metadata describing its source path and type.

- Markdown files yield a single full-document segment.
- Task JSON is walked recursively, but only a curated allowlist of paths is translated, including persona fields, task instructions, reason-for-call, natural-language assertions, and user-or assistant-visible message history.
- Database JSON/TOML is translated only at conservative natural-language leaf keys such as name, title, description, summary, and notes. Domain-specific additions may extend this set: the airline domain also translates `address1`, `address2`, and `city` (user profile address text that is natural language and safe to localize).
- Tool Python files are parsed with `ast`; only docstrings attached to `@is_tool` or `@is_discoverable_tool` methods are extracted.
- Google-style tool docstrings are decomposed into short description, long description, parameter descriptions, returns text, and raises text so that each part can be translated independently and later reassembled.
- Schema Python files are converted into JSON artifacts whose translatable fields include class descriptions, field descriptions, enum values, and `Literal[...]` alternatives.
- Tool return files expose both exact response strings and template strings as translatable segments.

Artifact Class	Source Files	Translatable Content	Runtime Alignment Purpose
Task setup	tasks.json	Persona text, instructions, reason-for-call, history	Defines user personas and evaluation-facing assertions.
Policy & workflow	*.md	Full document prose	Governs the agent’s allowed behavior and procedures.
Tool descriptions	tools.py, user_tools.py	Docstrings of tool-decorated methods	Conditions how the agent interprets tool semantics.
Schemas & literals	data_model.py, user_data_model.py	Descriptions, enum values, and Literal[...]	Defines choices appearing in dynamic tool schemas.
Database content	db.json, db.toml	Natural-language leaves (e.g., names, notes)	Supplies textual fields surfaced through tool results.
Tool return messages	tool_returns.json	Exact messages and parameterized templates	Ensures dynamic strings returned by tools remain localized.

Table 2: Artifact classes processed during offline translation and their role in maintaining runtime evaluation alignment.

Step 3. Canonical-token masking. Before translation, the pipeline collects and masks strings that must remain canonical: IDs such as `order_*` or `booking_*`, status values, structural task markers, tool names, and docstring section headers such as `Args` or `Returns`. Protection is partly global and partly contextual. For example, airline literals such as `basic_economy` or `round_trip` are protected when they occur in `cabin-class` or `trip-type` contexts, but not when similar surface strings appear as ordinary prose. The masking layer replaces protected strings with opaque placeholders and restores them after translation.

Step 4. Schema-first literal translation. Schema enum labels and `Literal[...]` values are translated first in a dedicated literal mode. This produces language-specific schema artifacts from which the pipeline builds a domain literal map, including alias forms such as underscore, space-separated, or hyphenated variants of the same canonical value. This step fixes the localized terminology before any longer-form prose is translated.

Step 5. Standard translation with glossary injection. All remaining segments, including task prose, policy text, tool docstrings, database leaves, and tool-return templates, are translated in standard mode. The model receives the schema-derived literal map as a glossary, and segments are pre-masked so that localized forms are restored consistently. Requests are deduplicated when possible, batched as structured JSON, and executed concurrently through LiteLLM. The current implementation requires the exact Vertex route `vertex_ai/gemini-3.1-flash-lite-preview`. Batch failures can be retried, recursively split, or

rerun individually if placeholder restoration fails.

Step 6. Format-aware writing and manifest recording. The translated text is written back using format-specific writers: markdown is emitted directly; JSON/TOML files are patched only at extracted addresses; tool docstrings are reconstructed and saved as `tools.json` or `user_tools.json`; schema artifacts are written as `data_model.json` and `user_data_model.json`; and DB files with no translated leaves are still copied through so the translated directory remains complete. Each language directory also receives `translation_manifest.json`, which records the output file, component, source language, target language, model, translation timestamp, and SHA-256 fingerprints of the source files. This manifest enables later staleness checks without forcing retranslation.

A.4 Phase 2. Runtime Localization

The offline artifacts cover static content, but the agent also sees dynamic tool schemas and tool outputs at inference time. These are handled by a runtime localization layer built from the paired source and translated schema artifacts.

Step 7. Build runtime localization maps. A `SchemaRuntimeLocalizer` constructs four resources from the source and localized schema artifacts: a description map, a canonical-to-localized literal map, a localized-to-canonical inverse map, and optional maps for exact and templated tool-return localizations.

Step 8. Localize the tool schema shown to the agent. The environment’s `get_tools()` path is

wrapped so that the agent receives localized tool schemas: descriptions are translated, enum choices are shown in the target language, and default or example literal values are localized where appropriate. The tool implementation itself is unchanged.

Step 9. Normalize localized arguments and localize tool outputs. Before tool execution, localized enum arguments supplied by the agent are mapped back to their canonical English values. After execution, structured response payloads and tool-return messages are localized back into the target language so the interaction remains monolingual from the agent’s perspective.

Step 10. Canonicalize localized values for evaluation. Prior to metric computation, localized payloads are canonicalized back to English, making pass-rate comparisons directly comparable across languages and against the original English benchmark.

A.5 Translated artifact statistics

Tables 3, 4, 5, and 6 provide summary statistics for translated artifacts resulting from offline translation phase (Appendix A.3).

Domain	Files / Language	Total Translated
Airline	5	25
Retail	5	25
Telecom	13	65
Total	23	115

Table 3: No. static artifact files translated per domain. Each domain is translated into 5 target languages.

Domain	Tasks	Docstrings	Policies	Return Msgs
Airline	250	14	5	0
Retail	570	16	5	0
Telecom	570	43	25	60

Table 4: No. instances per type translated across all 5 target languages.

Domain	Total Models	Value Sets	Localized Values
Airline	23	15	21
Retail	15	6	14
Telecom	18	12	51

Table 5: Schema artifacts and literal inventories.

Domain	Collections	Record Breakdown
Airline	3	flights: 300, users: 500, reservations: 2k
Retail	3	products: 50, users: 500, orders: 1k
Telecom	6	plans: 5, devices: 29, lines: 9, customers: 4, bills: 6

Table 6: Database artifacts. Structure is preserved; only designated leaf fields are translated.

B Detail on Mixed Language Tool Adaptation

For the mixed language tool adaptation, our goal is to investigate the stability of agents when handling tools in mixed languages. We conducted a similar experiment to the (S3) L2 Tool Adaptation where we modify the language of the tool specification into L2, but instead of converting to a specific L2, in this experiment we convert the tool specifications to a mix of several L2 languages. Specifically, we explore 5 different mixed language settings by gradually increasing the number of languages in the tool specification while maintaining a similar composition across domains: Mix-2 uses English and Thai, Mix-3 adds Vietnamese, Mix-4 adds Indonesian, and Mix-5 adds Chinese. For each run, the language of a tool is fixed for the entire run to reduce noise from random sampling. For instance, if `get_item` tool is assigned to Thai, it remains in Thai across all examples. For each non-English language added, three tools are localized into that language, while the remaining tools stay in English; thus, the English-tool count decreases as more languages are added.

C Hyperparameters

Hyperparameter	Agent LLM	User Sim. LLM
Temperature	0.0	0.0
Nucleus sampling (p)	1.0	1.0
Max. generation tokens	—	—

Table 7: Inference hyperparameters for the L2 Interaction Setting. Values follow the τ^2 -Bench defaults; parameters not listed are left at the provider model defaults.

All models in the L2 Interaction Setting are run with the default inference parameters from τ^2 -Bench, as shown in Table 7.

Scenario	Agent LLM Cost (USD)
English Baseline	\$70.40
L2 Tool Adaptation	\$340.73
L2 Interaction	\$403.27
L2 Domain Adaptation	\$440.43
Total	\$1,254.83

Table 8: Approximate agent LLM inference cost (USD) for each experimental scenario, aggregated across three evaluation trials and estimated from token usage with published API pricing.

D Compute Budget

For agent model inference, we use two models across all experimental settings: GPT-5-mini accessed via the OpenRouter API, and Kimi-K2.5 accessed via Azure AI Foundry. Table 8 summarizes the approximate total inference cost per scenario. For user model inference, we self-host Qwen3-235B-A22B-Instruct-2507 on 8xA100 GPUs (40GB each), resulting in a total of 4,032 GPU hours.

In the translation pipeline section, we used Gemini 3.1 Flash-Lite Preview for three domains and five target languages. Aggregated across five languages, the pipeline consumes 8.97M input tokens and 7.21M output tokens. Based on the Gemini 3.1 Flash-Lite pricing of \$0.25/M input tokens and \$1.50/M output tokens, the total translation cost is approximately \$13.05.

E Agent language use and drift analysis

We characterize agent language behavior along three axes: how often the agent produces text in the expected target language (Figure 9), where in a conversation any off-target text appears (Figure 10), and how much of the off-target text is specifically English (Figure 11).

Correctness is high overall and degrades only in crosslingual interaction. Figure 9 shows that L2 Tool is essentially perfect across all tested languages (mean turn correctness ≈ 1.00 for Thai, Vietnamese, Indonesian, Chinese, and Filipino), indicating that translating tool schemas does not by itself induce agent language drift: with an English dialogue the agent reliably stays in English. L2 Domain also remains high (0.92–0.98), with a gentle ordering in which Thai and Vietnamese are strongest (0.98) and Filipino is weakest (0.92). The clearest degradation is confined to L2 Interaction,




Figure 9: Average turn-level agent language correctness by setting and L2. Language correctness is calculated as the fraction of eligible turns whose detected language (by fastText) matches the expected L2 for one turn. We exclude simulated user language correctness because it is consistently above 0.95 across scenarios.



Figure 10: Proportion of non-L2 turns as a function of agent turn position in a conversation, separated by setting and by speaker role (agent vs. user).

where correctness falls to 0.75–0.82 and is again lowest for Filipino (0.75) and Chinese (0.78). The gap between settings is large relative to the gap between languages: moving from L2 Domain to L2 Interaction costs roughly 0.15–0.20 in correctness for every language, whereas the spread across languages within a setting is at most ~ 0.07 . Drift is therefore primarily a property of the interaction regime rather than of any single language.

Text of off-target language accumulates over the course of a conversation. Figure 10 decomposes this drift by turn position. In L2 Tool, the non-L2 proportion stays negligible throughout for both roles, consistent with scenario design where dialogue should be in English. In L2 Interaction and L2 Domain, the agent’s non-L2 share rises monotonically with turn index (Pearson $r \approx 0.65$ and $r \approx 0.73$ between turn position and off-target proportion, respectively), pointing to an accumulation or self-priming effect rather than a fixed per-turn error rate. The two settings differ in their dynamics: in L2 Interaction both the simulated user and the agent contribute off-target turns and the magnitude is largest (the agent share peaks near 0.12 in late turns), whereas in L2 Domain the simulated user




Figure 11: Per-task share of agent turns emitted in English, by target language, for the two scenarios in which the dialogue language is L2. Each dot is a task; the overlaid line tracks the per-language mean.

remains essentially on-target (≈ 0.015 throughout) and the modest rise to ≈ 0.07 is agent-initiated and concentrated in the later turns. This is consistent with Figure 9: L2 Interaction is where both parties’ code-switching reinforces one another.

The off-target text is predominantly English, and only for Filipino is it non-trivial. Figure 11 shows that when the agent leaves the target language it overwhelmingly switches to English, but that this is rare for most languages. For Thai, Vietnamese, Indonesian, and Chinese the per-task median English share is 0 in both settings, with right-skewed distributions whose means sit at only a few percent. Filipino is the clear outlier: its English share is highest in both settings (mean 0.11, median 0.06 in L2 Interaction; mean 0.08, median 0.04 in L2 Domain), and the per-language mean line rises toward Filipino in both panels. The right-skew indicates that English switching is driven by a minority of tasks rather than being uniform, so even the elevated Filipino mean reflects a tail of high-switching tasks rather than pervasive drift.

Drift is real but is not the dominant explanation for performance. Taken together, language drift

is genuine and scenario-dependent: strongest in L2 Interaction, concentrated in Filipino, accumulating over turns, and realized mainly as English code-switching. It is nonetheless too localized and too weakly associated with task success to serve as the primary explanation for performance differences. Across all scenarios, agent language correctness is essentially uncorrelated with pass^3 (overall Pearson $r = -0.10$, $R^2 \approx 0.01$, $n = 166$); the only setting with a strong positive correlation is L2 Tool, where correctness is saturated near 1.00 and thus carries no usable variance, while in L2 Interaction and L2 Domain the scenario-level correlations are weak ($|r| \lesssim 0.25$). Language drift is therefore best read as a measurable secondary symptom of the harder crosslingual regimes (L2 Interaction) rather than a direct cause of the pass^3 gaps reported in the main text.

F All results

Domain	Model	EN			
		p@1	p ²	p ³	ρ ³
Airline	GPT-5-mini	0.693	0.593	0.540	0.779
	Qwen3-235B	0.500	0.387	0.340	0.680
	Kimi-K2.5	0.707	0.600	0.540	0.764
Retail	GPT-5-mini	0.655	0.497	0.412	0.629
	Qwen3-235B	0.579	0.415	0.342	0.591
	Kimi-K2.5	0.561	0.442	0.395	0.704
Telecom	GPT-5-mini	0.713	0.588	0.509	0.714
	Qwen3-235B	0.357	0.205	0.132	0.370
	Kimi-K2.5	0.997	0.953	0.930	0.933

Table 9: English Baseline (S1) results. Where EN denotes English, p@1 denotes pass@1, p² denotes pass², and p³ denotes pass³. Bold values indicate the best score within each domain–metric group. Qwen3-235B refers to Qwen3-235B-A22B-Instruct-2507.

Domain	Model	VI				TH				ID				Zh				TL			
		p@1	p ²	p ³	ρ ³	p@1	p ²	p ³	ρ ³	p@1	p ²	p ³	ρ ³	p@1	p ²	p ³	ρ ³	p@1	p ²	p ³	ρ ³
Airline	GPT-5-mini	0.640	0.500	0.420	0.656	0.593	0.473	0.400	0.674	0.627	0.447	0.360	0.574	0.620	0.460	0.380	0.613	0.620	0.527	0.480	0.774
	Qwen3-235B	0.480	0.347	0.280	0.583	0.487	0.373	0.320	0.657	0.493	0.347	0.260	0.527	0.460	0.300	0.220	0.478	0.440	0.313	0.260	0.591
	Kimi-K2.5	0.733	0.660	0.600	0.818	0.687	0.580	0.520	0.757	0.740	0.653	0.600	0.811	0.760	0.640	0.580	0.763	0.747	0.653	0.600	0.804
Retail	GPT-5-mini	0.696	0.570	0.500	0.718	0.699	0.538	0.430	0.615	0.684	0.547	0.456	0.667	0.675	0.506	0.412	0.610	0.655	0.518	0.430	0.656
	Qwen3-235B	0.556	0.418	0.333	0.600	0.526	0.383	0.325	0.617	0.608	0.462	0.386	0.635	0.591	0.442	0.351	0.594	0.573	0.444	0.386	0.674
	Kimi-K2.5	0.687	0.550	0.465	0.677	0.573	0.430	0.360	0.627	0.640	0.483	0.386	0.603	0.722	0.597	0.526	0.729	0.675	0.541	0.474	0.701
Telecom	GPT-5-mini	0.561	0.430	0.368	0.656	0.605	0.436	0.342	0.565	0.582	0.453	0.404	0.693	0.728	0.608	0.518	0.711	0.740	0.617	0.535	0.723
	Qwen3-235B	0.260	0.129	0.061	0.236	0.222	0.108	0.061	0.276	0.333	0.170	0.105	0.316	0.228	0.132	0.097	0.423	0.155	0.088	0.061	0.396
	Kimi-K2.5	0.912	0.836	0.772	0.846	0.854	0.763	0.702	0.822	0.915	0.839	0.772	0.843	0.930	0.868	0.816	0.877	0.883	0.784	0.702	0.795

Table 10: L2 Interaction (S2) results across Vietnamese (VI), Thai (TH), Indonesian (ID), Chinese (ZH), and Filipino (TL). Here, p@1 denotes pass@1, p² denotes pass², and p³ denotes pass³. Bold values indicate the best score within each domain–metric group. Qwen3-235B refers to Qwen3-235B-A22B-Instruct-2507.

Domain	Model	VI				TH				ID				Zh				TL			
		p@1	p ²	p ³	ρ ³	p@1	p ²	p ³	ρ ³	p@1	p ²	p ³	ρ ³	p@1	p ²	p ³	ρ ³	p@1	p ²	p ³	ρ ³
Airline	GPT-5-mini	0.653	0.527	0.460	0.704	0.620	0.473	0.400	0.645	0.687	0.567	0.480	0.699	0.687	0.593	0.520	0.757	0.667	0.580	0.540	0.810
	Qwen3-235B	0.447	0.313	0.260	0.582	0.380	0.253	0.220	0.579	0.400	0.233	0.180	0.450	0.480	0.313	0.200	0.417	0.467	0.320	0.240	0.514
Retail	GPT-5-mini	0.681	0.532	0.430	0.631	0.646	0.477	0.377	0.584	0.675	0.526	0.430	0.637	0.719	0.579	0.482	0.670	0.661	0.515	0.439	0.664
	Qwen3-235B	0.547	0.412	0.351	0.642	0.561	0.421	0.351	0.626	0.550	0.418	0.351	0.638	0.594	0.433	0.360	0.606	0.570	0.421	0.342	0.600
Telecom	GPT-5-mini	0.664	0.550	0.482	0.726	0.713	0.588	0.518	0.727	0.713	0.596	0.518	0.727	0.705	0.576	0.500	0.709	0.719	0.594	0.509	0.708
	Qwen3-235B	0.254	0.135	0.088	0.346	0.339	0.208	0.140	0.413	0.269	0.155	0.105	0.390	0.222	0.135	0.096	0.432	0.301	0.167	0.105	0.349

Table 11: Tool Adaptation (S3) results in monolingual settings across Vietnamese (VI), Thai (TH), Indonesian (ID), Chinese (ZH), and Filipino (TL). Here, p@1 denotes pass@1, p² denotes pass², and p³ denotes pass³. Bold values indicate the best score within each domain–metric group. Qwen3-235B refers to Qwen3-235B-A22B-Instruct-2507.

Domain	Model	Bi				Tri				Quad				Multi			
		p@1	p ²	p ³	ρ ³	p@1	p ²	p ³	ρ ³	p@1	p ²	p ³	ρ ³	p@1	p ²	p ³	ρ ³
Airline	GPT-5-mini	0.680	0.560	0.480	0.706	0.660	0.560	0.480	0.727	0.640	0.520	0.480	0.750	0.640	0.513	0.440	0.688
	Qwen3-235B	0.440	0.280	0.200	0.455	0.413	0.267	0.200	0.484	0.413	0.287	0.240	0.581	0.413	0.267	0.220	0.533
Retail	GPT-5-mini	0.667	0.541	0.465	0.697	0.667	0.538	0.456	0.684	0.652	0.485	0.386	0.592	0.696	0.547	0.447	0.642
	Qwen3-235B	0.596	0.456	0.386	0.648	0.535	0.398	0.325	0.607	0.585	0.447	0.377	0.644	0.570	0.401	0.298	0.523
Telecom	GPT-5-mini	0.664	0.541	0.482	0.726	0.696	0.558	0.491	0.705	0.670	0.550	0.491	0.733	0.652	0.520	0.465	0.713
	Qwen3-235B	0.298	0.155	0.088	0.295	0.392	0.213	0.140	0.357	0.365	0.181	0.105	0.288	0.342	0.181	0.114	0.333

Table 12: Tool Adaptation (S3) results for multilingual settings, including bilingual (Bi), trilingual (Tri), quadlingual (Quad), and multilingual (Multi) configurations. Here, p@1 denotes pass@1, p² denotes pass², and p³ denotes pass³. Bold values indicate the best score within each domain–metric group. Qwen3-235B refers to Qwen3-235B-A22B-Instruct-2507.

Domain	Model	VI				TH				ID				Zh				TL			
		p@1	p ²	p ³	ρ^3	p@1	p ²	p ³	ρ^3	p@1	p ²	p ³	ρ^3	p@1	p ²	p ³	ρ^3	p@1	p ²	p ³	ρ^3
Airline	GPT-5-mini	0.540	0.400	0.320	0.593	0.593	0.453	0.360	0.607	0.600	0.540	0.500	0.833	0.607	0.493	0.440	0.725	0.560	0.433	0.380	0.679
	Qwen3-235B	0.487	0.460	0.440	0.903	0.460	0.413	0.380	0.826	0.480	0.440	0.420	0.875	0.500	0.473	0.460	0.920	0.473	0.420	0.400	0.846
	Kimi-K2.5	0.560	0.427	0.340	0.607	0.547	0.387	0.300	0.548	0.600	0.487	0.420	0.700	0.660	0.507	0.440	0.667	0.600	0.487	0.400	0.667
Retail	GPT-5-mini	0.585	0.427	0.333	0.569	0.325	0.175	0.123	0.378	0.541	0.392	0.325	0.601	0.596	0.450	0.368	0.617	0.488	0.333	0.263	0.539
	Qwen3-235B	0.289	0.205	0.149	0.516	0.178	0.143	0.132	0.742	0.243	0.170	0.140	0.576	0.266	0.181	0.149	0.560	0.260	0.187	0.158	0.608
	Kimi-K2.5	0.567	0.392	0.268	0.473	0.327	0.190	0.114	0.349	0.433	0.287	0.202	0.467	0.607	0.450	0.366	0.603	0.444	0.292	0.219	0.493
Telecom	GPT-5-mini	0.564	0.427	0.360	0.638	0.599	0.447	0.360	0.601	0.497	0.325	0.246	0.495	0.760	0.652	0.588	0.774	0.681	0.547	0.474	0.696
	Qwen3-235B	0.383	0.213	0.149	0.389	0.284	0.143	0.079	0.278	0.427	0.295	0.211	0.494	0.330	0.205	0.149	0.452	0.330	0.205	0.149	0.452
	Kimi-K2.5	0.699	0.529	0.404	0.578	0.693	0.523	0.43	0.620	0.798	0.681	0.588	0.737	0.863	0.763	0.684	0.793	0.743	0.582	0.474	0.638

Table 13: L2 Domain (S4) results across Vietnamese (VI), Thai (TH), Indonesian (ID), Chinese (ZH), and Filipino (TL). Here, p@1 denotes pass@1, p² denotes pass², and p³ denotes pass³. Bold values indicate the best score within each domain–metric group. Qwen3-235B refers to Qwen3-235B-A22B-Instruct-2507.

G Detailed Error Analysis Results

Setting	Variant	Agent Errors			User Errors		
		Critical	Minor	Correct	Critical	Minor	Correct
English Baseline	English	81	22	47	38	48	64
L2 Interaction	Vietnamese	93	25	32	43	48	59
	Thai	105	14	31	42	45	63
	Indonesian	88	18	44	36	49	65
	Chinese	90	18	42	44	51	55
	Filipino	94	18	38	53	36	61
Tool Adaptation	VI Tools	81	23	46	43	51	56
	TH Tools	88	28	34	39	45	66
	ID Tools	94	20	36	41	41	68
	ZH Tools	84	23	43	35	44	71
	FIL Tools	92	17	41	45	45	60
	Mix-2 (EN+TH)	90	24	36	52	38	60
	Mix-3 (+VI)	86	17	47	33	54	63
	Mix-4 (+ID)	89	26	35	35	45	70
	Mix-5 (+ZH)	102	15	33	40	41	69
L2 Domain	Vietnamese	83	27	40	43	41	66
	Thai	97	17	36	55	30	65
	Indonesian	92	24	34	42	39	69
	Chinese	103	15	32	40	48	62
	Filipino	101	17	32	50	47	53

Table 14: Error analysis on the **Airline** domain (gpt-5-mini agent, qwen3-235b user simulator, DeepSeek-V4-Flash judge, $n=150$ per condition). Counts of simulations by maximum severity level.

Setting	Variant	Agent Errors			User Errors		
		Critical	Minor	Correct	Critical	Minor	Correct
English Baseline	English	98	70	174	79	135	128
L2 Interaction	Vietnamese	109	78	155	96	103	143
	Thai	125	73	144	94	124	124
	Indonesian	117	65	160	86	121	135
	Chinese	109	73	160	81	126	135
	Filipino	103	88	151	105	129	108
Tool Adaptation	VI Tools	97	75	170	75	130	137
	TH Tools	102	79	161	70	123	149
	ID Tools	94	86	162	92	130	120
	ZH Tools	95	64	183	87	121	134
	FIL Tools	88	79	175	79	129	134
	Mix-2 (EN+TH)	81	81	180	58	132	152
	Mix-3 (+VI)	100	69	173	88	123	131
	Mix-4 (+ID)	100	72	170	88	112	142
	Mix-5 (+ZH)	87	93	162	82	129	131
L2 Domain	Vietnamese	112	93	137	90	122	130
	Thai	180	60	102	127	91	124
	Indonesian	124	76	142	93	122	127
	Chinese	131	84	127	101	113	128
	Filipino	131	74	137	123	118	101

Table 15: Error analysis on the **Retail** domain (gpt-5-mini agent, qwen3-235b user simulator, DeepSeek-V4-Flash judge, $n=342$ per condition). Counts of simulations by maximum severity level.

Setting	Variant	Agent Errors			User Errors		
		Critical	Minor	Correct	Critical	Minor	Correct
English Baseline	English	110	113	119	50	130	162
L2 Interaction	Vietnamese	101	116	125	52	146	144
	Thai	125	90	127	59	133	150
	Indonesian	122	100	120	60	136	146
	Chinese	95	76	171	60	129	153
	Filipino	91	88	163	62	141	139
Tool Adaptation	VI Tools	103	117	122	45	156	141
	TH Tools	87	106	149	46	147	149
	ID Tools	89	108	145	48	140	154
	ZH Tools	109	91	142	55	135	152
	FIL Tools	75	112	155	45	115	182
	Mix-2 (EN+TH)	91	106	145	54	138	150
	Mix-3 (+VI)	95	117	130	55	144	143
	Mix-4 (+ID)	91	102	149	56	128	158
	Mix-5 (+ZH)	94	101	147	43	148	151
L2 Domain	Vietnamese	125	71	146	54	112	176
	Thai	133	85	124	64	105	173
	Indonesian	139	81	122	59	120	163
	Chinese	103	86	153	45	128	169
	Filipino	131	89	122	64	130	148

Table 16: Error analysis on the **Telecom** domain (gpt-5-mini agent, qwen3-235b user simulator, DeepSeek-V4-Flash judge, $n=342$ per condition). Counts of simulations by maximum severity level.

	Eng. Baseline		L2 Interaction										Tool Adaptation										L2 Domain					
Error Tag			TH	TH	TH	TH	TH	TH	TH	TH	TH	ZH	ZH	Mix-2	Mix-3	Mix-3	Mix-5	TH	TH	TH	TH	TH	TH					
Guideline Violation	C100-M2	C116-M38	C113-M28	C89-M31	C85-M27	C117-M108	C82-M19	C105-M30	C87-M32	C92-M26	C89-M21	C110-M22	C97-M16	C96-M24	C103-M26	C90-M36	C130-M28	C115-M18	C108-M21	C114-M24								
Hallucination	C86-M7	C67-M2	C46-M8	C48-M5	C99-M6	C83-M5	C90-M4	C55-M2	C55-M6	C42-M4	C46-M5	C33-M2	C35-M5	C25-M5	C7-M2	C82-M2	C14-M10	C35-M5	C73-M8	C86-M7								
Unethical Behavior	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1	C77-M1								
Incorrect Interpretation	C33-M25	C37-M20	C56-M30	C47-M33	C46-M38	C26-M36	C25-M29	C46-M35	C33-M40	C24-M31	C39-M29	C34-M29	C33-M24	C36-M34	C46-M21	C34-M37	C71-M35	C43-M25	C47-M32	C86-M30								
Interpretation Error	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0	C0-M0								
Irrelevant Tool Call	C0-M0	C3-M4	C0-M4	C0-M2	C0-M4	C0-M4	C0-M0	C0-M0	C0-M0	C0-M0	C0-M1	C0-M1	C1-M3	C0-M2	C0-M2	C1-M0	C0-M0	C0-M0	C0-M0	C0-M1								
Missed Required Action	C11-M16	C82-M28	C98-M18	C74-M15	C85-M20	C90-M25	C79-M22	C75-M22	C78-M17	C75-M27	C85-M22	C78-M20	C78-M14	C87-M33	C84-M21	C77-M15	C86-M16	C96-M25	C94-M19	C110-M23								
Tool Error	C0-M1	C2-M4	C0-M1	C0-M1	C1-M3	C0-M0	C0-M2	C0-M1	C2-M1	C0-M0	C0-M0	C0-M1	C2-M0	C0-M1	C0-M1	C0-M1	C1-M0	C0-M1	C0-M1	C0-M4								
Generative Terminology	C0-M1	C0-M1	C18-M6	C0-M1	C2-M1	C21-M4	C19-M0	C19-M0	C19-M0	C19-M0	C19-M0	C19-M0	C19-M0	C25-M1	C26-M1	C25-M1	C25-M1	C25-M1	C25-M1	C27-M1								
Revealed Info Early	C0-M1	C0-M2	C0-M1	C0-M1	C0-M1	C2-M2	C3-M1	C1-M1	C0-M0	C0-M1	C1-M1	C1-M1	C2-M1	C0-M0	C0-M0	C1-M2	C0-M0	C0-M1	C1-M2	C0-M0								
Tool Call Argument Error	C5-M4	C3-M1	C7-M6	C4-M6	C3-M7	C1-M4	C2-M6	C2-M6	C1-M4	C1-M0	C5-M0	C3-M3	C2-M5	C1-M4	C1-M2	C5-M3	C5-M3	C7-M8	C6-M4	C6-M2								
Tool Call Error	C0-M1	C2-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1	C0-M1								
Wrong Sequence	C9-M3	C9-M7	C8-M3	C3-M4	C11-M4	C7-M3	C12-M6	C3-M3	C5-M3	C12-M6	C5-M2	C5-M3	C1-M1	C1-M1	C7-M1	C8-M3	C7-M3	C13-M5	C12-M2	C11-M5								

Table 17: Agent error tag counts — all settings, **Airline** domain (gpt-5-mini agent, qwen3-235b user simulator, DeepSeek-V4-Flash judge, $n=150$ per condition). **C**: = critical-severity count; **M**: = minor-severity count. Each value is the number of simulations in which the tag appeared at that severity level at least once.

	Eng_Baseline						L2 Interaction								Tool Adaptation							L2 Domain					
Error Tag	CIN	V1	TH	IDH	ZH	FIL	V1	TH	ID	ZH	FIL	Mix-1	Mix-2	Mix-3	Mix-4	Mix-5	V1	TH	ID	ZH	FIL	CIN	V1	TH	IDH	ZH	FIL
Guideline Violation	E13 M26	C18 M22	C19 M25	C15 S30	C28 M23	C26 M31	C19 M36	C9 M17	C16 M13	C14 H25	C22 M18	C23 M19	C16 M22	C19 M20	C18 M54	C16 M20	C29 M14	C20 M13	C18 B19	C21 M24							
Hallucination	C26 M27	C22 M24	C31 M24	C19 M37	C23 M31	C39 M25	C39 M28	C39 M20	C25 M21	C32 M26	C35 M29	C40 M19	C30 M26	C25 M21	C31 M20	C33 M20	C35 M22	C48 M13	C34 M13	C33 B18	C33 M18						
Inconsistent Tool Behavior	C23 M24	C19 M21	C24 M22	C20 M23	C30 M21	C39 M25	C39 M28	C39 M20	C25 M21	C32 M26	C35 M29	C40 M19	C30 M26	C25 M21	C31 M20	C33 M20	C35 M22	C48 M13	C34 M13	C33 B18	C33 M18						
Incorrect Interpretation	C4 M27	C3 M18	C2 M59	C2 M54	C3 M39	C1 M17	C2 M12	C1 M31	C3 M2	C3 M5	C1 M4	C6 M8	C4 M8	C3 M4	C1 M7	C3 M10	C3 M2	C3 M2	C3 M6	C10 M13	C5 M8						
Interruption Error	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0						
Invalid Input	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0						
Missed Required Action	C0 M1	C3 M1	C2 M0	C0 M1	C1 M2	C2 M1	C1 M1	C1 M0	C3 M1	C1 M0	C1 M2	C1 M1	C0 M1	C1 M3	C0 M1	C2 M3	C0 M1	C0 M0	C0 M0	C0 M0	C0 M0						
Other	C0 M1	C0 M1	C0 M0	C0 M0	C0 M2	C1 M0	C0 M0	C0 M2	C0 M0	C1 M0	C0 M1	C0 M0	C0 M3	C1 M0	C0 M0	C0 M2	C0 M0	C0 M1	C0 M0	C0 M0	C0 M0						
Tool Call Termination	C0 M1	C12 M1	C14 M7	C0 M0	C0 M1	C0 M1	C0 M1	C0 M1	C12 M1	C0 M1	C0 M7	C0 M1	C0 M1	C2 M1	C0 M1	C0 M1	C0 M1	C0 M1	C0 M1	C0 M1	C0 M1						
Revealed Info Error	C0 M2	C0 M0	C0 M2	C0 M0	C0 M1	C0 M2	C0 M2	C0 M2	C0 M0	C1 M1	C0 M0	C0 M1	C0 M0	C0 M0	C0 M1	C0 M0	C0 M1	C0 M0	C0 M0	C0 M0	C0 M0						
Tool Call Argument Error	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0						
Wrong Context	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0						

Table 18: User error tag counts — all settings, **Airline** domain (gpt-5-mini agent, qwen3-235b user simulator, DeepSeek-V4-Flash judge, $n=150$ per condition). **C**: = critical-severity count; **M**: = minor-severity count. Each value is the number of simulations in which the tag appeared at that severity level at least once.

Eng. Baseline			L2 Interaction								Tool Adaptation								L2 Domain							
Error Tag	TH	ID	TH	ID	TH	ID	TH	ID	TH	ID	TH	ID	TH	ID	TH	ID	TH	ID	TH	ID	TH	ID				
Guideline Violation	C70-M74	C99-M80	C105-M81	C108-M83	C71-M29	C86-M94	C68-M52	C70-M166	C69-M69	C62-M55	C78-M56	C41-M55	C67-M48	C63-M54	C78-M72	C73-M76	C128-M54	C92-M53	C112-M93	C92-M59						
Hallicination	C74-M11	C59-M12	C15-M51	C17-M17	C13-M11	C10-M10	C9-M11	C8-M13	C3-M11	C16-M67	C15-M10	C11-M12	C17-M18	C12-M18	C18-M17	C30-M16	C18-M17	C9-M18	C14-M13							
Incorrect Behavior	C74-M11	C59-M12	C15-M51	C17-M17	C13-M11	C10-M10	C9-M11	C8-M13	C3-M11	C16-M67	C15-M10	C11-M12	C17-M18	C12-M18	C18-M17	C30-M16	C18-M17	C9-M18	C14-M13							
Incorrect Interpretation	C34-M39	C32-M54	C42-M59	C38-M54	C23-M64	C37-M39	C31-M31	C40-M47	C28-M39	C40-M34	C36-M34	C22-M33	C29-M32	C38-M43	C27-M65	C48-M88	C142-M62	C50-M56	C63-M55							
Interpation Error	C70-M50	C60-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0	C40-M0							
Tool Call Error	C1-M53	C1-M3	C1-M4	C1-M5	C1-M5	C1-M3	C1-M5	C1-M5	C1-M5	C1-M5	C1-M5	C1-M5	C1-M5	C1-M5	C1-M5	C1-M5	C1-M5	C1-M5	C1-M5							
Missed Required Action	C75-M53	C78-M62	C55-M51	C58-M55	C93-M53	C94-M63	C62-M59	C79-M46	C74-M64	C69-M65	C58-M56	C63-M71	C73-M49	C59-M67	C73-M66	C160-M62	C75-M65	C98-M83	C101-M81							
Other	C0-M2	C0-M3	C1-M7	C2-M1	C0-M3	C0-M3	C0-M3	C0-M3	C0-M4	C2-M1	C1-M2	C1-M2	C0-M4	C0-M1	C0-M6	C0-M3	C1-M5	C1-M2	C1-M5							
Incorrect Termination	C0-M2	C0-M3	C1-M7	C2-M1	C0-M3	C0-M3	C0-M3	C0-M3	C0-M4	C2-M1	C1-M2	C1-M2	C0-M4	C0-M1	C0-M6	C0-M3	C1-M5	C1-M2	C1-M5							
Revealed Info Error	C0-M0	C1-M5	C4-M4	C2-M1	C1-M2	C2-M3	C2-M3	C1-M2	C1-M5	C0-M6	C0-M0	C3-M1	C0-M5	C1-M1	C1-M3	C2-M3	C2-M3	C2-M3	C2-M2							
Tool Call Argument Error	C8-M57	C67-M18	C8-M58	C9-M10	C5-M5	C9-M17	C5-M7	C5-M7	C8-M9	C8-M3	C8-M6	C5-M14	C7-M4	C8-M9	C5-M8	C5-M49	C48-M39	C39-M28	C42-M37	C73-M34						
Tool Call Error	C20-M6	C24-M22	C20-M17	C19-M20	C19-M16	C23-M15	C16-M14	C15-M21	C18-M18	C3-M13	C11-M9	C5-M14	C7-M10	C11-M13	C14-M11	C24-M8	C19-M21	C12-M15	C13-M12							

Table 19: Agent error tag counts — all settings, **Retail** domain (gpt-5-mini agent, qwen3-235b user simulator, DeepSeek-V4-Flash judge, $n=342$ per condition). **C**: = critical-severity count; **M**: = minor-severity count. Each value is the number of simulations in which the tag appeared at that severity level at least once.

[illegible]

Table 20: User error tag counts — all settings, **Retail** domain (gpt-5-mini agent, qwen3-235b user simulator, DeepSeek-V4-Flash judge, $n=342$ per condition). **C**: = critical-severity count; **M**: = minor-severity count. Each value is the number of simulations in which the tag appeared at that severity level at least once.

Error Tag	Eng. Baseline		L2 Intervention					Tool Adaptation										L2 Domain				
	EN		VI		TH		ID		VI		TH		ID		VI		TH		ID		VI	
	FIL		FIL		FIL		FIL		FIL		FIL		FIL		FIL		FIL		FIL		FIL	
Guideline Violation	C38 M81	C77 M128	C99 M105	C103 M116	C43 M54	C61 M68	C71 M90	C49 M77	C57 M94	C77 M61	C49 M88	C80 M83	C78 M81	C70 M75	C72 M60	C117 M104	C106 M97	C98 M91	C78 M81	C106 M115		
Hallucination	C4 M38	C6 M126	C7 M915	C6 M114	C73 M54	C4 M11	C4 M18	C4 M18	C4 M18	C23 M14	C2 M18	C4 M15	C3 M14	C4 M17	C3 M14	C7 M16	C5 M11	C5 M19	C5 M3	C3 M11		
Model Error	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94		
Incorrect Interpretation	C39 M94	C38 M84	C47 M100	C43 M100	C39 M63	C29 M60	C41 M30	C17 M91	C30 M30	C37 M74	C26 M68	C26 M77	C36 M98	C26 M96	C40 M80	C30 M69	C54 M93	C53 M77	C29 M71	C54 M85		
Interpretation Error	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0	C0 M0		
Model Error	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94	C39 M94		
Model Required Action	C17 M74	C106 M132	C142 M118	C138 M132	C109 M107	C107 M97	C104 M154	C38 M124	C104 M133	C134 M142	C56 M147	C107 M153	C100 M189	C106 M135	C118 M118	C151 M96	C166 M176	C178 M108	C111 M91	C168 M111		
Model Error	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6	C1 M6		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12	C18 M1	C25 M1	C28 M1	C28 M1	C28 M3	C28 M3	C28 M3	C28 M4	C33 M0	C137 M12	C23 M0	C30 M3	C39 M2	C23 M3	C37 M5		
Model Error	C25 M4	C25 M6	C36 M24	C28 M2	C32 M12																	

Table 21: Agent error tag counts — all settings, **Telecom** domain (gpt-5-mini agent, qwen3-235b user simulator, DeepSeek-V4-Flash judge, $n=342$ per condition). **C**: = critical-severity count; **M**: = minor-severity count. Each value is the number of simulations in which the tag appeared at that severity level at least once.

Error Tag	Eng. Baseline	L2 Interaction					Tool Adaptation										L2 Domain				
	EN	VI	TH	ID	ZH	FIL	VI	TH	ID	ZH	FIL	Mix-2	Mix-3	Mix-4	Mix-5	VI	TH	ID	ZH	FIL	
Guideline Violation	C:25 M:42	C:10 M:93	C:30 M:77	C:30 M:53	C:30 M:82	C:54 M:105	C:33 M:52	C:36 M:61	C:16 M:74	C:19 M:64	C:35 M:58	C:15 M:49	C:13 M:52	C:15 M:77	C:12 M:56	C:13 M:42	C:17 M:46	C:11 M:59	C:23 M:85	C:22 M:85	
Hallucination	C:52 M:85	C:34 M:96	C:47 M:67	C:68 M:74	C:52 M:85	C:29 M:68	C:45 M:91	C:48 M:88	C:51 M:78	C:49 M:55	C:32 M:60	C:48 M:93	C:40 M:95	C:38 M:69	C:74 M:83	C:39 M:48	C:49 M:56	C:57 M:54	C:37 M:61	C:52 M:47	
Inconsistent Behavior	C:24 M:194	C:19 M:207	C:26 M:186	C:33 M:139	C:40 M:197	C:52 M:194	C:17 M:186	C:47 M:175	C:15 M:180	C:22 M:160	C:23 M:135	C:28 M:158	C:32 M:209	C:23 M:169	C:14 M:184	C:18 M:146	C:16 M:151	C:20 M:140	C:10 M:115	C:20 M:142	
Incorrect Interpretation	C:1 M:50	C:6 M:30	C:2 M:36	C:8 M:49	C:7 M:38	C:3 M:31	C:4 M:30	C:6 M:36	C:5 M:49	C:6 M:52	C:4 M:46	C:4 M:38	C:7 M:35	C:4 M:39	C:4 M:29	C:3 M:30	C:7 M:34	C:8 M:25	C:6 M:28	C:7 M:42	
Interruption Error	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	C:0 M:0	
Irrelevant Tool Call	C:0 M:0	C:0 M:1	C:0 M:1	C:0 M:4	C:0 M:0	C:0 M:0	C:0 M:1	C:0 M:1	C:0 M:8	C:0 M:3	C:0 M:10	C:0 M:0	C:0 M:1	C:0 M:3	C:0 M:0	C:0 M:1	C:0 M:0	C:0 M:1	C:0 M:3	C:0 M:3	
Missed Required Action	C:3 M:24	C:3 M:9	C:6 M:16	C:1 M:22	C:3 M:12	C:2 M:29	C:2 M:11	C:2 M:30	C:2 M:21	C:5 M:16	C:9 M:11	C:3 M:22	C:2 M:8	C:4 M:17	C:1 M:21	C:6 M:8	C:2 M:12	C:1 M:9	C:3 M:14	C:2 M:9	
Other	C:1 M:10	C:1 M:9	C:0 M:13	C:1 M:9	C:0 M:28	C:0 M:15	C:0 M:15	C:1 M:3	C:0 M:12	C:0 M:32	C:0 M:11	C:1 M:9	C:0 M:16	C:1 M:10	C:1 M:10	C:1 M:11	C:0 M:33	C:0 M:20	C:1 M:8	C:0 M:27	
Premature Termination	C:13 M:16	C:9 M:24	C:12 M:20	C:10 M:18	C:14 M:21	C:17 M:16	C:5 M:23	C:3 M:14	C:12 M:20	C:8 M:15	C:11 M:14	C:11 M:15	C:12 M:17	C:15 M:14	C:5 M:16	C:11 M:23	C:17 M:12	C:14 M:28	C:9 M:18	C:11 M:31	
Revealed Info Early	C:0 M:0	C:0 M:0	C:0 M:2	C:0 M:1	C:1 M:3	C:0 M:3	C:0 M:1	C:0 M:0	C:0 M:4	C:0 M:15	C:1 M:1	C:1 M:1	C:0 M:0	C:0 M:0	C:0 M:0	C:1 M:0	C:1 M:1	C:1 M:2	C:0 M:2	C:0 M:1	
Tool Call Argument Error	C:0 M:3	C:2 M:7	C:0 M:2	C:2 M:1	C:1 M:0	C:1 M:2	C:0 M:10	C:0 M:5	C:1 M:9	C:0 M:1	C:0 M:2	C:1 M:4	C:4 M:8	C:0 M:1	C:0 M:1	C:0 M:5	C:0 M:2	C:0 M:1	C:0 M:3	C:0 M:2	
Tool Call Schema Error	C:0 M:1	C:5 M:6	C:0 M:0	C:0 M:2	C:0 M:8	C:0 M:5	C:0 M:3	C:0 M:6	C:0 M:5	C:0 M:2	C:0 M:4	C:0 M:3	C:0 M:1	C:0 M:5	C:0 M:1	C:0 M:0	C:0 M:1	C:0 M:22	C:1 M:1	C:0 M:2	
Wrong Sequence	C:0 M:0	C:0 M:0	C:0 M:11	C:4 M:6	C:2 M:5	C:0 M:4	C:0 M:7	C:1 M:2	C:0 M:6	C:2 M:3	C:0 M:3	C:0 M:10	C:0 M:1	C:0 M:5	C:0 M:2	C:0 M:2	C:1 M:1	C:0 M:1	C:0 M:1	C:0 M:1	

Table 22: User error tag counts — all settings, **Telecom** domain (gpt-5-mini agent, qwen3-235b user simulator, DeepSeek-V4-Flash judge, $n=342$ per condition). **C:** = critical-severity count; **M:** = minor-severity count. Each value is the number of simulations in which the tag appeared at that severity level at least once.